


(a) Throughput vs Context Length



(b) DBA Speedup vs Baseline

