

*Baseline  
collapses*

## Single-Token Decode Throughput vs Context Length

